

fstat(2) - Linux man page

Name

stat, fstat, lstat - get file status

Synopsis

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <unistd.h>
```

```
int stat(const char *path, struct stat *buf);
```

```
int fstat(int fd, struct stat *buf);
```

```
int lstat(const char *path, struct stat *buf);
```

Feature Test Macro Requirements for glibc (see [feature_test_macros\(7\)](#)):

lstat():

```
_BSD_SOURCE || _XOPEN_SOURCE >= 500 || _XOPEN_SOURCE &&  
_XOPEN_SOURCE_EXTENDED  
|| /* Since glibc 2.10: */ _POSIX_C_SOURCE >= 200112L
```

Description

These functions return information about a file. No permissions are required on the file itself, but-in the case of **stat()** and **lstat()** - execute (search) permission is required on all of the directories in *path* that lead to the file.

stat() stats the file pointed to by *path* and fills in *buf*.

lstat() is identical to **stat()**, except that if *path* is a symbolic link, then the link itself is stat-ed, not the file that it refers to.

fstat() is identical to **stat()**, except that the file to be stat-ed is specified by the file descriptor *fd*.

All of these system calls return a *stat* structure, which contains the following fields:

```
struct stat {  
    dev_t      st_dev;      /* ID of device containing file */  
    ino_t      st_ino;      /* inode number */  
    mode_t     st_mode;     /* protection */  
    nlink_t    st_nlink;    /* number of hard links */  
    uid_t      st_uid;      /* user ID of owner */  
    gid_t      st_gid;      /* group ID of owner */  
    dev_t      st_rdev;     /* device ID (if special file) */
```

```

off_t      st_size;      /* total size, in bytes */
blksize_t  st_blksize;   /* blocksize for file system I/O */
blkcnt_t   st_blocks;    /* number of 512B blocks allocated */
time_t     st_atime;     /* time of last access */
time_t     st_mtime;     /* time of last modification */
time_t     st_ctime;     /* time of last status change */
};

```

The `st_dev` field describes the device on which this file resides. (The **major**(3) and **minor**(3) macros may be useful to decompose the device ID in this field.)

The `st_rdev` field describes the device that this file (inode) represents.

The `st_size` field gives the size of the file (if it is a regular file or a symbolic link) in bytes. The size of a symbolic link is the length of the pathname it contains, without a terminating null byte.

The `st_blocks` field indicates the number of blocks allocated to the file, 512-byte units. (This may be smaller than `st_size/512` when the file has holes.)

The `st_blksize` field gives the "preferred" blocksize for efficient file system I/O. (Writing to a file in smaller chunks may cause an inefficient read-modify-rewrite.)

Not all of the Linux file systems implement all of the time fields. Some file system types allow mounting in such a way that file and/or directory accesses do not cause an update of the `st_atime` field. (See *noatime*, *nodiratime*, and *relatime* in **mount**(8), and related information in **mount**(2).) In addition, `st_atime` is not updated if a file is opened with the **O_NOATIME**; see **open**(2).

The field `st_atime` is changed by file accesses, for example, by **execve**(2), **mknod**(2), **pipe**(2), **utime**(2) and **read**(2) (of more than zero bytes). Other routines, like **mmap**(2), may or may not update `st_atime`.

The field `st_mtime` is changed by file modifications, for example, by **mknod**(2), **truncate**(2), **utime**(2) and **write**(2) (of more than zero bytes). Moreover, `st_mtime` of a directory is changed by the creation or deletion of files in that directory. The `st_mtime` field is *not* changed for changes in owner, group, hard link count, or mode.

The field `st_ctime` is changed by writing or by setting inode information (i.e., owner, group, link count, mode, etc.).

The following POSIX macros are defined to check the file type using the `st_mode` field:

S_ISREG(m)

is it a regular file?

S_ISDIR(m)

directory?

S_ISCHR(m)

character device?

S_ISBLK(m)

block device?

S_ISFIFO(m)

FIFO (named pipe)?

S_ISLNK(m)

symbolic link? (Not in POSIX.1-1996.)

S_ISSOCK(m)

socket? (Not in POSIX.1-1996.)

The following flags are defined for the *st_mode* field:

The set-group-ID bit (**S_ISGID**) has several special uses. For a directory it indicates that BSD semantics is to be used for that directory: files created there inherit their group ID from the directory, not from the effective group ID of the creating process, and directories created there will also get the **S_ISGID** bit set. For a file that does not have the group execution bit (**S_IXGRP**) set, the set-group-ID bit indicates mandatory file/record locking.

The sticky bit (**S_ISVTX**) on a directory means that a file in that directory can be renamed or deleted only by the owner of the file, by the owner of the directory, and by a privileged process.

Return Value

On success, zero is returned. On error, -1 is returned, and *errno* is set appropriately.

Errors

EACCES

Search permission is denied for one of the directories in the path prefix of *path*. (See also [path_resolution](#)(7).)

EBADF

fd is bad.

EFAULT

Bad address.

ELOOP

Too many symbolic links encountered while traversing the path.

ENAMETOOLONG

path is too long.

ENOENT

A component of *path* does not exist, or *path* is an empty string.

ENOMEM

Out of memory (i.e., kernel memory).

ENOTDIR

A component of the path prefix of *path* is not a directory.

EOVERFLOW

path or *fd* refers to a file whose size, inode number, or number of blocks cannot be represented in, respectively, the types *off_t*, *ino_t*, or *blkcnt_t*. This error can occur when, for example, an application compiled on a 32-bit platform without `-D_FILE_OFFSET_BITS=64` calls **stat()** on a file whose size exceeds $(1 \ll 31) - 1$ bytes.

Conforming To

These system calls conform to SVr4, 4.3BSD, POSIX.1-2001.

According to POSIX.1-2001, **lstat()** on a symbolic link need return valid information only in the *st_size* field and the file-type component of the *st_mode* field of the *stat* structure. POSIX.-2008 tightens the specification, requiring **lstat()** to return valid information in all fields except the permission bits in *st_mode*.

Use of the *st_blocks* and *st_blksize* fields may be less portable. (They were introduced in BSD. The interpretation differs between systems, and possibly on a single system when NFS mounts are involved.) If you need to obtain the definition of the *blkcnt_t* or *blksize_t* types from `<sys/stat.h>`, then define **_XOPEN_SOURCE** with the value 500 or greater (before including *any* header files).

POSIX.1-1990 did not describe the **S_IFMT**, **S_IFSOCK**, **S_IFLNK**, **S_IFREG**, **S_IFBLK**, **S_IFDIR**, **S_IFCHR**, **S_IFIFO**, **S_ISVTX** constants, but instead demanded the use of the macros **S_ISDIR()**, etc. The **S_IF*** constants are present in POSIX.1-2001 and later.

The **S_ISLNK()** and **S_ISSOCK()** macros are not in POSIX.1-1996, but both are present in POSIX.1-2001; the former is from SVID 4, the latter from SUSv2.

UNIX V7 (and later systems) had **S_IREAD**, **S_IWRITE**, **S_IEXEC**, where POSIX prescribes the synonyms **S_IRUSR**, **S_IWUSR**, **S_IXUSR**.

Other systems

Values that have been (or are) in use on various systems:

A sticky command appeared in Version 32V AT&T UNIX.

Notes

Since kernel 2.5.48, the *stat* structure supports nanosecond resolution for the three file timestamp fields. Glibc exposes the nanosecond component of each field using names of the form *st_atim.tv_nsec* if the **_BSD_SOURCE** or **_SVID_SOURCE** feature test macro is defined. These fields are specified in POSIX.1-2008, and, starting with version 2.12, glibc also exposes these field names if **_POSIX_C_SOURCE** is defined with the value 200809L or greater, or **_XOPEN_SOURCE** is defined with the value 700 or greater. If none of the aforementioned macros are defined, then the nanosecond values are exposed with names of the form *st_atimensec*. On file systems that do not support subsecond timestamps, the nanosecond fields are returned with the value 0.

On Linux, **lstat()** will generally not trigger automounter action, whereas **stat()** will (but see **fstatat(2)**).

For most files under the */proc* directory, **stat()** does not return the file size in the *st_size* field; instead the field is returned with the value 0.

Underlying kernel interface

Over time, increases in the size of the *stat* structure have led to three successive versions of **stat()**: *sys_stat()* (slot *__NR_oldstat*), *sys_newstat()* (slot *__NR_stat*), and *sys_stat64()* (new in kernel 2.4; slot *__NR_stat64*). The glibc **stat()** wrapper function hides these details from applications, invoking the most recent version of the system call provided by the kernel, and repacking the returned information if required for old binaries. Similar remarks apply for **fstat()** and **lstat()**.

Example

The following program calls **stat()** and displays selected fields in the returned *stat* structure.

```
#include <sys/types.h>
#include <sys/stat.h>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>

int
main(int argc, char *argv[])
{
```

```
struct stat sb;

if (argc != 2) {
    fprintf(stderr, "Usage: %s <pathname>\n", argv[0]);
    exit(EXIT_FAILURE);
}

if (stat(argv[1], &sb) == -1) {
    perror("stat");
    exit(EXIT_FAILURE);
}

printf("File type:                ");

switch (sb.st_mode & S_IFMT) {
    case S_IFBLK:  printf("block device\n");          break;
    case S_IFCHR:  printf("character device\n");       break;
    case S_IFDIR:  printf("directory\n");              break;
    case S_IFIFO:  printf("FIFO/pipe\n");              break;
    case S_IFLNK:  printf("symlink\n");                break;
    case S_IFREG:  printf("regular file\n");           break;
    case S_IFSOCK: printf("socket\n");                 break;
    default:       printf("unknown?\n");               break;
}

printf("I-node number:           %ld\n", (long) sb.st_ino);

printf("Mode:                     %lo (octal)\n",
       (unsigned long) sb.st_mode);

printf("Link count:                %ld\n", (long) sb.st_nlink);
printf("Ownership:                 UID=%ld   GID=%ld\n",
       (long) sb.st_uid, (long) sb.st_gid);

printf("Preferred I/O block size: %ld bytes\n",
       (long) sb.st_blksize);
printf("File size:                  %lld bytes\n",
       (long long) sb.st_size);
printf("Blocks allocated:           %lld\n",
       (long long) sb.st_blocks);

printf("Last status change:         %s", ctime(&sb.st_ctime));
printf("Last file access:          %s", ctime(&sb.st_atime));
printf("Last file modification:     %s", ctime(&sb.st_mtime));
```

```
    exit(EXIT_SUCCESS);  
}
```

See Also

[access\(2\)](#), [chmod\(2\)](#), [chown\(2\)](#), [fstatat\(2\)](#), [readlink\(2\)](#), [utime\(2\)](#), [capabilities\(7\)](#), [symlink\(7\)](#)

Referenced By

[dirfd\(3\)](#), [explain\(1\)](#), [explain_fstat\(3\)](#), [explain_fstat_or_die\(3\)](#), [fseek\(3\)](#), [isatty\(3\)](#), [lam_rfpix\(2\)](#), [obsolete\(2\)](#), [pidfile\(3\)](#), [progress\(1\)](#), [shm_open\(3\)](#), [shm_overview\(7\)](#), [syscalls\(2\)](#), [ttyname\(3\)](#), [xfs_io\(8\)](#)